

# AI assignment - 2

**SEARCH & OPTIMIZATION**

Atreya M - CS22B009

Bhukya Lachiram Nayak - CS22B012

# Frozen Lake - Environment

## In Frozen Lake

- Each cell is a state.
- Actions = {Up, Down, Left, Right}.
- Transitions are deterministic (non-slippery).
- Goal is to find the **shortest path to the goal**.



# 1. BRANCH & BOUND

Branch and Bound is a best-first search algorithm that systematically explores paths by expanding the most promising nodes first—those with the lowest estimated total cost ( $\text{cost\_so\_far} + \text{heuristic}$ ).

## How it Works

1. **Start from initial state (S).**
2. Maintain a **priority queue** (min-heap) of paths, prioritized by their cost + heuristic.
3. At each step:
  - Pop the **path with the lowest priority**.
  - If it reaches the **goal (G)**, return it.
  - Otherwise, **expand** the path by moving in all possible directions (up/down/left/right).
  - Add each new path to the queue if the resulting state hasn't been visited.
4. **Repeat** until goal is found or queue is empty

$\text{cost\_so\_far}$ : Number of steps taken.

heuristic: Estimated cost to goal (we use **Manhattan distance**).


## Advantages

- Explores only the best candidate paths first.
- Finds **optimal** solution if heuristic is admissible .

## Disadvantages

- May explore many unnecessary paths if heuristic is weak.
- Memory intensive (stores many partial paths).

BnB



# Iterative Deepening A\*

IDA\* combines **depth-first search (DFS)** with the cost-based pruning of **A\***. It performs a series of **depth-limited searches**, increasing the cost bound at each iteration.

## HOW IT WORKS:

- **Set an initial cost bound** using:  $\text{bound} = \text{heuristic}(\text{start})$
- Run a **DFS**, only exploring paths where:  $\text{cost\_so\_far} + \text{heuristic} \leq \text{bound}$
- If the goal is found, return it
- If not found, **increase the bound** to the minimum value that exceeded the last one and **repeat**.

## Process

- Each iteration is a DFS up to a given cost threshold.
- The threshold is incrementally increased (hence "iterative deepening").

## Advantages

- **Memory efficient** (like DFS).
- **Optimal solution** .
- Works well in large state spaces.

## Disadvantages

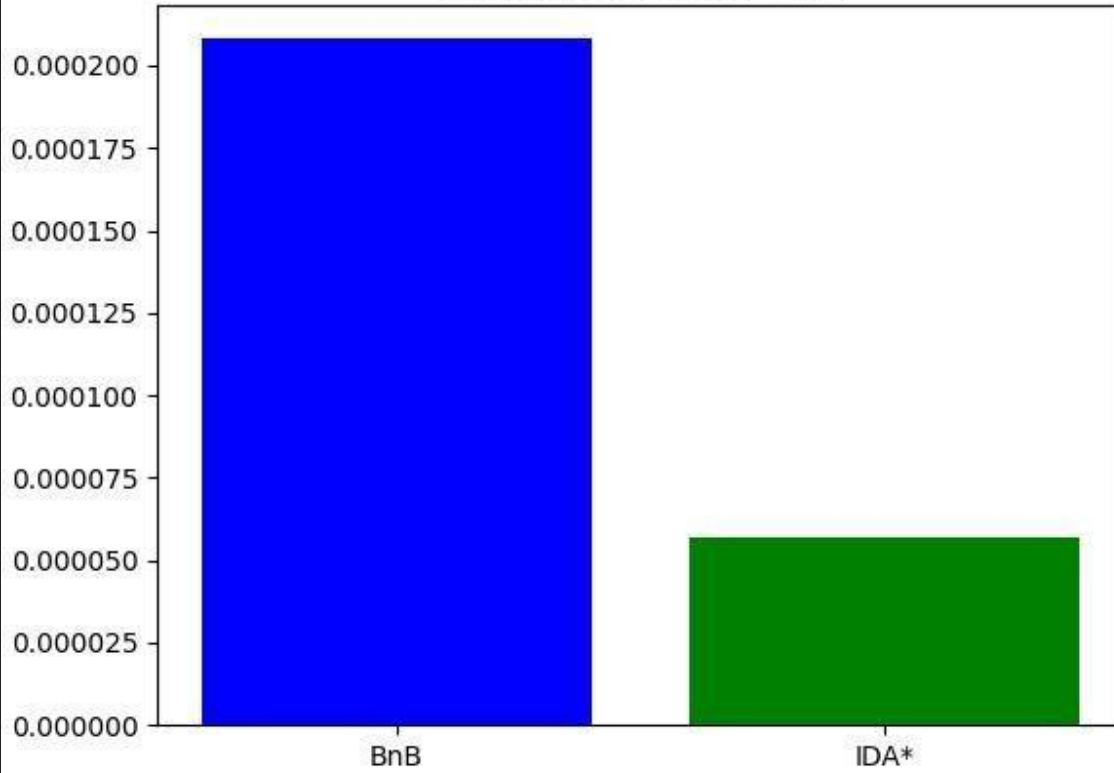
- **Slower than A\*** due to repeated work.
- Requires careful tuning to avoid recomputation overhead.

IDA\*





Avg Time to Reach Goal



## What's the key difference with BnB and IDA\*?

**BnB** tries to explore all possibilities but skips branches that are worse than the best solution found so far.

**IDA\*** combines the benefits of A\* (heuristics) with iterative deepening to explore efficiently.

Even with pruning, BnB often has to **explore a large part of the search space**, especially early on. If a better solution is found late, earlier pruning isn't very effective.

# Traveling Salesman Problem - Environment

## Experiment Setup

- **Environment:** 20 random 2D points (`mk_env`).
- **Objective:** Minimize total path length (TSP).
- **Evaluation:**
  - Each algorithm runs **5 times**.
  - Average cost is calculated and visualized.
  - Final tour is saved as a plot (`hc_tour.png`, `sa_tour.png`).

Generating the cities using `np.random.rand()` , and we can compute distances between cities using numpy op

### 3.HILL CLIMBING

Hill Climbing is a greedy local search algorithm. It continuously moves towards a better solution by exploring its neighbors, hoping to reach the optimal solution.

#### How It Works

1. Start with a random tour (a random permutation of points).
2. Repeatedly:
  - Pick two indices  $i < j$ .
  - Reverse the segment between  $i$  and  $j$  (a 2-opt move).
  - If the new tour has a lower cost, accept it.
  - Else, revert the change.
3. Run this for a fixed time ( $t$  seconds in the code).

Cost Function: The total distance of the tour

## **Advantages of Hill Climbing**

1. Simple to implement
2. Fast and efficient
3. Uses little memory

## **Disadvantages of Hill Climbing**

1. Can get stuck in local optima
2. No backtracking
3. Not good for complex landscapes

# Simulated Annealing

Simulated Annealing is an advanced local search algorithm inspired by the annealing process in metallurgy. Unlike hill climbing, SA accepts worse solutions with a certain probability to escape local minima.

## How It Works

1. Start with a **random solution** .
2. Initialize **temperature**  $T$
3. Repeat until time expires or  $T$  is very small:
  - Choose two points  $i < j$  and reverse the segment (2-opt move).
  - Compute **cost difference**  $d = \text{new\_cost} - \text{old\_cost}$ .
  - Accept the move if:
    - $d < 0$  (better solution), or
    - $\exp(-d/T) > \text{random}()$  (probabilistic acceptance).
  - Gradually **decrease temperature**:  $T *= \text{cool\_rate}$  (like 0.995).

## Temperature Schedule

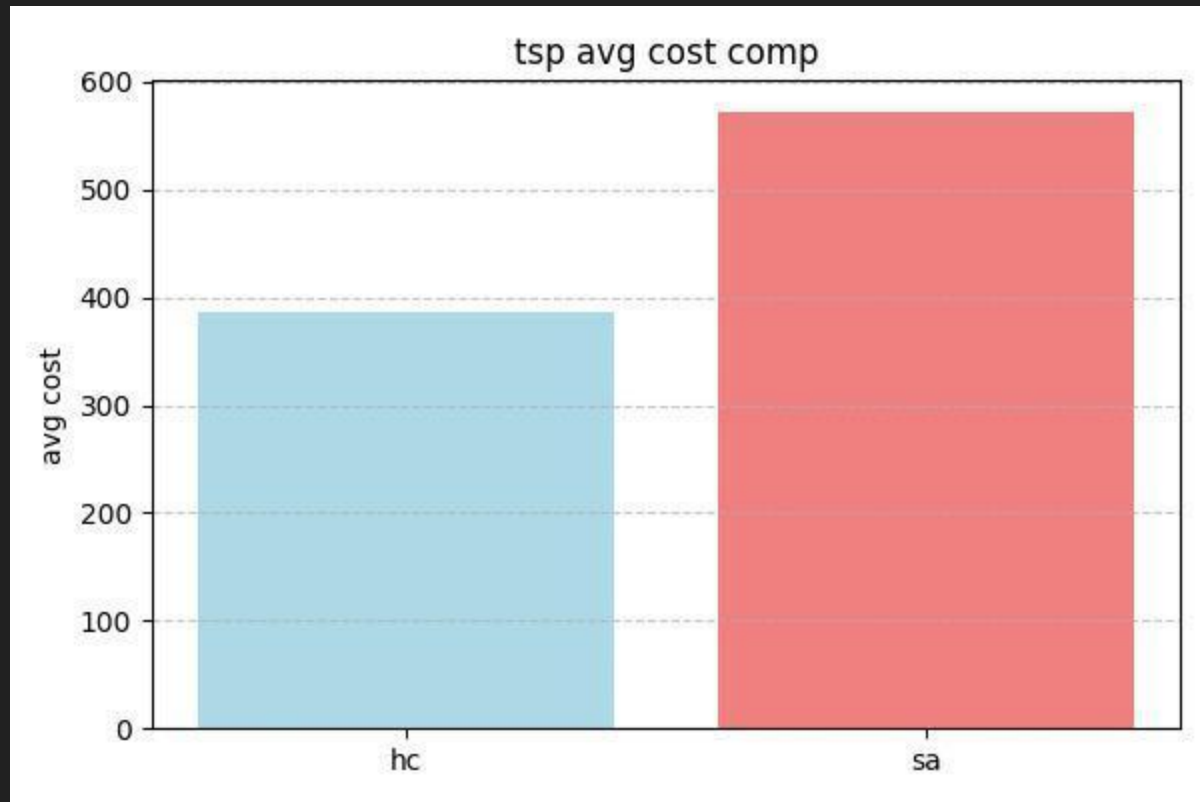
- Starts high to allow exploration (accepting worse paths).
- Gets stricter over time (low temperature = greedy).

## Advantages

- Escapes local minima.
- Can find **global optima** over time.

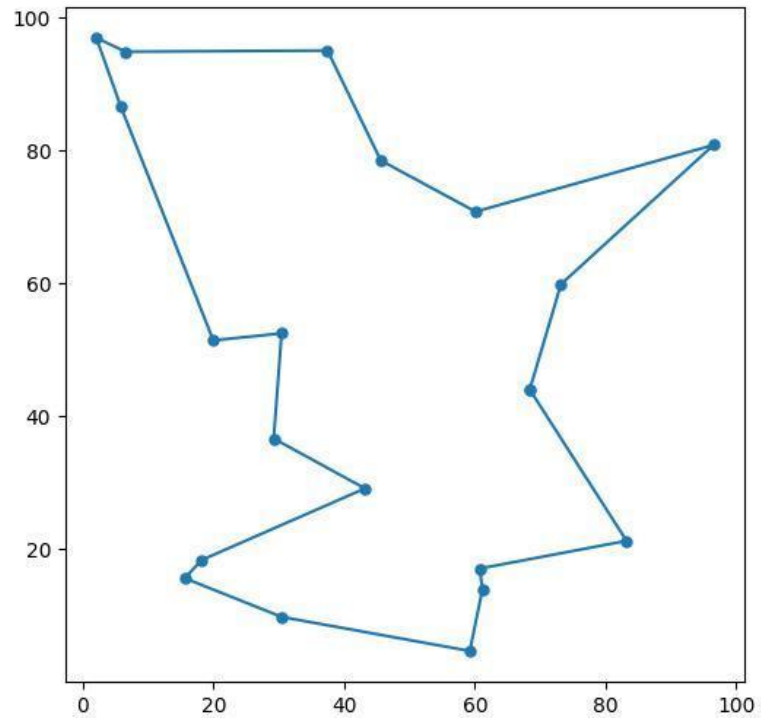
## Disadvantages

- Slower than hill climbing.
- Requires careful tuning (initial temperature, cooling rate).





hc tour



sa tour

